

USER GUIDE · OMEGA COMPUTE

Compute Site Pro Forma

Screen a building's existing electrical service for GPU compute pods beside its EV chargers, a battery and its own load. Size it, choose the deal and present the after-tax IRR in your company's brand.

Workspace
Compute Site Pro Forma
Saved scenarios...

Guide
Save scenario
you@yourcompany.com

- Site & service
TX - 253 kW limit
- Building load
Typical multifamily load (ass...
- EV charging
12 ports - managed
- Compute
192 H200 GPUs
- Load balance
ADVANCE - peak 211 kW
- Deal & financing
Own the power and pod laye...
- Results & report**
31.3% IRR

STEP 7 OF 7

Results & report

The same returns, IRR build and investor deck as the BESS Pro Forma, from the same finance engine.

Headline returns

Own the power and pod layer for a revenue share · 100% equity, unlevered

\$1.41M
total investment sought
100% equity

31.29%
after-tax unlevered IRR
over the 10-year term

\$2,968,212
total investor returns
over the life of the project

~2.7 years
after-tax payback
cash + depreciation + ITC

IRR build, after tax

Cash only (pre-tax operating cash flow)	36.3%
Depreciation effect	(5.7) pts
ITC effect	+0.7 pts
Total after-tax IRR	31.3%

Each step is rounded to a tenth of a point, as the deck prints it.

More metrics

NPV at 10%	\$1,360,732
Year-one distribution	\$561,564
Multiple on invested capital	3.10x
Pre-tax IRR	36.27%
Installed cost	\$1,407,240

LIVE SUMMARY

31.29%
after-tax unlevered IRR · 10-year term

Investment sought	\$1.41M
After-tax payback	~2.7 years
NPV @ 10%	\$1.36M
Year-one distribution	\$562K
Peak at the meter	211 kW of 253 kW

Site screen [ADVANCE](#)

Up to date · compute-site-1

What it answers

Can this building's existing electrical service carry GPU compute pods beside its EV chargers, a battery and its own load, and what does that return? The Compute Site Pro Forma answers both from one hourly model of the site. It then prices three ways to own the site and prints the investor deck.

This is the metro-edge model: inference pods placed on power a building already has, often the power built for its EV chargers. The model runs every hour of a year on that one service. The building comes first, then the compute sold under contract, then the charging. The battery and the on-demand GPUs share what is left. Nothing is added up in your browser: every figure comes back from ClearSky's server.

Where to find it

- **Omega Workspace** > **Finance** hexagon > **Compute Site Pro Forma**, right after the VPP Earnings Simulator. It is also in the **Compute** hexagon.
- **Site Map** (the editor) > **Compute** tab > **Size** > **Site Screen** screens the drawing and opens the pro forma with the site filled in. The last page explains it.
- Direct: silmarillion.clearskyomega.com/compute-proforma. Add **?example=1** to open it on the example.
- It belongs to **Omega Compute**. Field and Pro plans and trials include it; on a package, add Omega Compute from the Modules page.

The fastest start: run the example

Press **Run the example**. It is on step 1, in the Live summary on the right, and on the empty Load balance and Results. The example fills every step and runs at once. If you had typed anything, **Put mine back** in the notice at the bottom of the screen restores it.

The example	An illustrative apartment building in Austin, ZIP 78701. Replace any of it with your own.
Service	400 A at 480 V, three-phase, carrying up to 253 kW continuously
Compute	192 H200 GPUs in 4 pods
Charging	12 ports at 11.5 kW, managed, residents charging overnight
Battery	100 kW / 200 kWh
Deal	Own the power layer
What it returns	ADVANCE on the site screen, a peak of 211 kW of 253 kW at the meter, and a 31.29% after-tax IRR on \$1.41M.

The service is tight on purpose: the twelve chargers at full power would overrun it, so managed charging, the battery and the on-demand GPUs all have work to do.

STEP 7 OF 7

Results & report

The same returns, IRR build and investor deck as the BESS Pro Forma, from the same finance engine.

Nothing has run yet. Run the example for a finished pro forma, or enter your site from step 1.

Run the example

Start at step 1

LIVE SUMMARY

after-tax IRR

Enter the ZIP and the service size to run the model.

Run the example

Before anything has run, the Results and the Live summary offer the example and step 1.

How the page works

Seven steps down the left, the Live summary on the right. Every change runs the model again on the server after a short pause. On a phone the steps are Back and Next at the bottom, and the summary is the strip under the header.

1. **Site & service:** where, and what the panel can carry.
2. **Building load:** an 8760, bills or a typical shape, and the tariff.
3. **EV charging:** ports, sessions and when cars come.
4. **Compute:** GPUs, who buys the hours, prices over time.
5. **Load balance:** the battery, the hourly fit and the sizing sweep.
6. **Deal & financing:** three ways to own it, tax and debt.
7. **Results & report:** the IRR, the investor deck and saved scenarios.

Once the model has run, a line under each step's name recaps it. A step with a problem is ringed in red. **Save scenario** at the top keeps the whole set; **Saved scenarios** reopens one.

A grey figure in a field is what a blank field runs on. Type over it to change it.

LIVE SUMMARY

31.29%

after-tax unlevered IRR · 10-year term

Investment sought **\$1.41M**

After-tax payback **~2.7 years**

NPV @ 10% **\$1.36M**

Year-one distribution **\$562K**

Peak at the meter **211 kW of 253 kW**

Site screen

ADVANCE

Up to date · compute-site-1

The Live summary.

1 Site & service

The **ZIP** is the one field the model cannot run without. It sets the market, the climate the typical load follows and the state's tax rate. **Market** reads the ZIP unless you pick one. **Host type** (Multifamily, Office, Retail, Industrial) sets the building's load shape and rate class. The names and the address go on the deck.

The site

[Run the example](#)

Scenario / project name

Example: 4 GPU pods on a 400 A apartment service (illustra

Sponsor

Who is raising the capital

Host

Example apartments

Host type

Multifamily

Office

Retail

Industrial

Sets the building's load shape and rate class.

Street

City

Austin

ZIP *

78701

Utility

e.g. Austin Energy

Market

From the ZIP

The site. The name, sponsor and host print on the investor deck.

The **electrical service** is the limit everything else is held under. Enter the rating from the main breaker on the panel schedule, in amps with the voltage and phases, or its capacity in kW. The model may load it to the **continuous loading limit**, 80% by default as the NEC sizes continuous load. A planned **service upgrade** adds capacity and its cost, capitalised as 15-year property.

The electrical service

From the panel schedule or the utility

Service rating

400

A

Voltage

480 V

Phases

Three-phase

...or its capacity in kW

from the amps

kW

Continuous loading limit

80

%

Power factor

0.95

= 316 kW from 400 A at 480 V

NEC sizes continuous load at 80% of the rating.

[> A service upgrade](#)

The service. 400 A at 480 V, three-phase, carrying up to 253 kW continuously.

2 Building load & tariff

The building's own load always comes first on the service. Give the model the best record you have:

- **8760 / interval:** an hourly or 15-minute file, in kWh or kW. The best.
- **Bills (12–24 mo):** each month's kWh and dollars. Three or more months with dollars calibrate the rate to what the building pays.

- **Typical:** a typical shape for the host type, scaled to the annual kWh if you have it. The results say it was assumed.

Typical

Bills (12–24 mo)

8760 / interval

Annual kWh

typical for the host type

kWh

The building load. The example uses the typical shape for an apartment building.

The **tariff** is optional. Enter the on-peak and off-peak energy prices, the demand charge and the on-peak hours, or paste an OpenEI URDB tariff record. Without one, the market's planning rate is used and a warning says so.

Tariff

Optional

Without it the market's planning rate is used, calibrated to the bills' dollars when three or more months carry them.

On-peak energy

Off-peak energy

Demand charge

\$

/kWh

\$

/kWh

\$

/kW-mo

On-peak starts

On-peak ends

:00

:00

[> An OpenEI URDB tariff record](#)

The tariff. Blank fields run on the market's planning rate.

3 EV charging

Charging is placed after the building and the contracted compute, inside what the service has left.

- **Charging ports and power per port:** Level 2 is typically 7.2, 11.5 or 19.2 kW; DC fast 62.5 or 150 kW.
- **Managed charging** (on by default) moves each session into its cheapest hours without setting a new peak. Off, cars charge at full power on arrival, held under the service by the panel's load management (NEC 625.42).
- **Who charges, and when:** residents overnight, workplace, retail and public, a fleet depot, or a custom arrival window. Sessions per port, weekend sessions, energy per session and the stay follow the pattern; type over any of them.
- Energy a car cannot get before it leaves is reported as **not delivered**, never assumed away.

Charging ports
 ports

Power per port
 kW
L2 7.2 · 11.5 · 19.2 kW; DC fast 62.5 · 150 kW.

Scheduling
☒ **Managed charging**
Off: cars charge on arrival, held under the service by the panel's load management (NEC 625.42).

Who charges, and when

Residents overnight
Workplace
Retail / public
Fleet depot
Custom window

Sessions per port
 a weekday

Weekend sessions
 × weekday

Energy per session
 kWh

Stay
 hours

The chargers. The example has twelve managed 11.5 kW ports for residents charging overnight.

Charging money holds the driver price and its escalator, the installed cost per port, the network fee, O&M, charger efficiency and payment processing. **Chargers already installed** takes the charger capital out of the model.

4 Compute

Choose the **GPU class** and the number of **pods**. The page shows each class's planning figures: RTX PRO 6000 at 0.60 kW a GPU, H200 at 0.70 kW and B300 at 1.30 kW, with 48, 48 and 32 GPUs a pod. Custom takes your own. Server overhead, cooling (PUE), availability and the serving and idle draw turn a GPU's rating into its draw at the meter.

The hardware

NVIDIA H200: 0.7 kW a GPU, up to 48 a pod

GPU class

RTX PRO 6000
H200
B300
Custom

Pods
 pods

GPUs per pod
 GPUs

Power per GPU
 kW

Server overhead
 %

Cooling (PUE)

Availability
 %

Draw serving
 % of rating

Draw idle
 % of rating

The hardware. 192 H200 GPUs in 4 pods.

Who buys the compute splits the GPUs across three contracts, each with a share and a price per GPU-hour:

- **Long-term offtake:** take-or-pay. These GPUs run every hour and are firm.
- **Local edge contracts:** firm inside their contract hours, at their average use.
- **On-demand overflow:** the one load the model may hold back to keep the peak down, and only where that earns more than the GPU-hours.

Who buys the compute

Use Laitent's published mix (40 / 30 / 30)

Contract	Share of GPUs		Price		Used on average
Long-term offtake	40	%	\$ 2.25	/GPU-hr	Take-or-pay
Local edge contracts	30	%	\$ 3.25	/GPU-hr	70 %
On-demand overflow	30	%	\$ 2.5	/GPU-hr	40 %

GPU rental indices, September 2026: H200 on-demand median \$4.40/GPU-hr (listings \$2.09–\$13.78), spot from \$1.99; B300 reserved median \$5.71, on-demand median \$7.87; RTX PRO 6000 from \$0.93, median \$1.48 (getdeploying.com GPU price index; jarvislabs.ai; thundercompute.com). The tier prices here are what an edge operator realises, below cloud list: a planning figure the operator's contract replaces.

Who buys the compute. Use Laitent's published mix sets 40 / 30 / 30.

The compute schedule decides when on-demand GPUs run: follow the site's peak (the default), always on, or off-peak hours only. **Prices over time** sets the yearly GPU-hour price decline, the offtake term, the first-year ramp, the GPU refresh and its cost, and the installed and running costs. The GPU prices are planning figures set below cloud list and marked as such: the compute operator's contract replaces them.

5 Battery & load balance

Enter the **battery** power and energy (blank means none) and any **solar**. Battery details hold the round trip, usable depth, fade, wear, O&M and replacement. The battery holds each month's peak where that pays, and keeps energy back for any hour the contracted load alone would overrun the service.

Battery and solar

Leave blank for none

Battery power

100 kW

Battery energy

200 kWh

Battery installed cost

\$ 550 /kWh

Solar

0 kW-dc

Solar installed cost

\$ 2.40 /W


> Battery details

Battery and solar. 100 kW / 200 kWh.

The site screen

The banner is the verdict, with the next gate after it.

ADVANCE	Fits the existing service. Above 85% of the limit it reads <i>tight margin</i> and the next gate is a load study.
VERIFY	More than 5% of the charging is not delivered, or every charger at full power would exceed the service: install load management (an EVEMS).
HOLD	The building and the contracted compute alone exceed the service, even with the battery.




The verdict and what the service carries. The bar is the peak at the meter against the limit; the orange mark is every charger at full power plus the firm load.

The four tiles are the headroom at the peak hour, the firm peak (the building and the contracted compute), the charging energy not delivered, and the on-demand GPU-hours held back.

Hour by hour


The bars stack every load in each hour: the building, the contracted compute, the on-demand compute, EV charging and battery charging. The black line is the meter. Where the bars rise above the line, the battery is covering the difference. The tabs show the peak summer weekday, the peak winter weekday and the busiest weekend day. Hover any hour for its figures, or open **The hours as a table**.

Hour by hour: Peak summer weekday (Jul 24)




The peak summer weekday. Charging runs overnight. In the evening peak the stack rises above the meter line: the battery is discharging to hold the month's peak down.

Monthly peak



Load duration

Every hour, sorted



Monthly peak (grey: without the battery or holding back GPUs; blue: at the meter) and every hour of the year sorted.

Year one on this service, under the charts, lists the GPU-hours and compute revenue, the charging delivered and its revenue, the battery's cycles and its effect on the bill, and the project's own electricity.

Size it

Find the best size runs pods against battery sizes, each cell a full year of hourly dispatch and a full pro forma. The best is the highest NPV that fits: no lost contracted compute and no more than 5% of the charging undelivered.

Shading follows NPV and a striped cell does not fit. **Use this size**, or a click on any cell that fits, applies it and runs the model again.

Size it

[Find the best size](#)

Pods against battery sizes, each a full year of hourly dispatch and the finance engine's after-tax return. The best is the highest NPV that fits: no lost firm compute and under 5% unserved charging.

Best: 5 pods (240 GPUs), no battery — 35.0% after-tax IRR, NPV \$1,828,277 on \$1,558,440 (owning the power layer).

[Use this size](#)

Pods	No battery	25 kW / 50 kWh	75 kW / 150 kWh	100 kW / 200 kWh	100 kW / 400 kWh
1	23.1% NPV \$279K peak 131 kW	21.7% NPV \$258K peak 118 kW	19.1% NPV \$216K peak 114 kW	17.9% NPV \$195K peak 113 kW	14.0% NPV \$111K peak 112 kW
2	29.4% NPV \$667K peak 166 kW	28.3% NPV \$646K peak 152 kW	26.3% NPV \$604K peak 146 kW	25.4% NPV \$583K peak 145 kW	22.1% NPV \$499K peak 143 kW
3	32.3% NPV \$1.06M peak 201 kW	31.5% NPV \$1.04M peak 188 kW	29.8% NPV \$993K peak 179 kW	29.1% NPV \$972K peak 179 kW	26.3% NPV \$888K peak 176 kW
4	34.0% NPV \$1.45M peak 236 kW	33.3% NPV \$1.42M peak 223 kW	31.9% NPV \$1.38M peak 213 kW	31.3% NPV \$1.36M peak 211 kW	28.9% NPV \$1.28M peak 210 kW
5	35.0% NPV \$1.83M peak 253 kW	34.5% NPV \$1.81M peak 253 kW	33.3% NPV \$1.77M peak 246 kW	32.8% NPV \$1.75M peak 245 kW	30.7% NPV \$1.66M peak 242 kW
6	33.5% 870 h over peak 253 kW	33.2% 335 h over peak 253 kW	32.4% 212 h over peak 253 kW	31.9% 212 h over peak 253 kW	30.1% 212 h over peak 253 kW

Each cell is a full year of hourly dispatch and the finance engine's after-tax return; the best is the highest NPV that fits the service with no lost firm compute and under 5% unserved charging. A striped cell does not fit; the outlined one is the best; the dark-edged one is the size above. Shading follows NPV. Each pod serves 39 kW at the meter; the limit is 253 kW.

The sizing sweep for the example: pods down the side, battery sizes across the top. Best: 5 pods (240 GPUs), no battery — 35.0% after-tax IRR, NPV \$1,828,277 on \$1,558,440 (owning the power layer).

Here the sweep prefers no battery: at this site's planning rate the battery holds the peak but saves less than it costs. That is what the sweep is for, finding it out before the capital is spent.

6 Deal & financing

The same hourly model is priced three ways, side by side. The one you choose gets the full pro forma.

- **Own and operate:** the project owns the GPUs, pods, chargers and battery, and sells compute through an operator's platform.
- **Own the power layer:** the investor owns the pods, power and chargers but not the GPUs, for a share of gross compute revenue (20% by default).

- **Lease the power:** the host leases the power to a compute operator for fixed rent with the electricity paid. No capital, so it shows the NPV of the rent and the property value it adds, never an IRR.

Ownership structure

Laitent's published host models, generalised

Own and operate

The project owns the GPUs, pods, chargers and battery; sells compute through an operator's platform and charging to drivers.

0.0% IRR · NPV -\$3.52M
\$9.70M capital · \$1.91M year-one EBITDA

Own the power layer

The investor owns the pods, power and chargers, not the GPUs, for a share of gross compute revenue.

31.3% IRR · NPV \$1.36M
\$1.41M capital · \$562K year-one EBITDA

Lease the power

The host leases the power to a compute operator for fixed rent, electricity paid.

No capital · NPV \$166K
\$25K year one · +\$417K property value

Share of gross compute

20%

Compute electricity

☒ The operator pays for the compute's electricity

Development and soft costs

8% of hard cost

Insurance

0.5% of capital

The three structures, each with its IRR, NPV, capital and year-one EBITDA. The example owns the power layer.

Term and rates holds the analysis term, discount rate, utility escalator, inflation and how the GPU refresh is paid. **Construction begins** and **In service** are a month and a year each: they decide which tax credits apply. Left blank, the credits are assumed kept and a warning says so.

Term and rates

Analysis term

10 years

Discount rate

10%

Utility escalator

3%/yr

Inflation

2.5%

GPU refresh paid

From that year's cash

Construction begins

Nov

2026

Sets which credits apply; blank assumes they all do.

In service

Jun

2027

The first operating month; the credit is claimed that year.

Term and rates, with the construction and in-service months the example filled in.

Tax & incentives runs through the BESS Pro Forma's tax engine: federal and state rates (a blank state rate uses the site's state), bonus depreciation, the tax appetite, whether the credit is claimed or sold (\$6418) and at what price, and

the battery credit's adders. **Financing** is 100% equity or with project debt, sized on loan-to-cost and coverage.

Tax & incentives

The BESS Pro Forma's tax engine

Federal rate

21

%

State rate

0

%

Bonus depreciation

0

%

100% is available for property acquired after 19 January 2025; use it only with the appetite to absorb it.

Tax appetite

Used as they arise

Carried forward (NOL)

An investor with other income uses the losses the year they arise; a stand-alone owner carries them forward.

The credit is

Claimed

Sold (\$6418)

Transfer price

\$

0.90

per \$1

When the credit is sold.

Storage credit

☒ Prevailing wage met
 ☐ Energy community

☐ Domestic content
 ☐ Meets the FEOC thresholds

The battery's investment tax credit and its adders.

Tax & incentives: the rates, the tax appetite, a credit claimed or sold, and the battery credit's adders.

7 Results & report

The same returns, IRR build and investor deck as the BESS Pro Forma, from the same finance engine.

Headline returns

Own the power and pod layer for a revenue share · 100% equity, unlevered

\$1.41M

total investment sought
100% equity

31.29%

after-tax unlevered IRR
over the 10-year term

\$2,968,212

total investor returns
over the life of the project

~2.7 years

after-tax payback
cash + depreciation + ITC

Headline returns: the investment sought, the after-tax IRR over the term, total investor returns and the after-tax payback.

ClearSky Energy Solutions · silmarillion.clearskyomega.com · page 12 of 15

IRR build, after tax

Cash only (pre-tax operating cash flow)	36.3%
Depreciation effect	(5.7) pts
ITC effect	+0.7 pts
Total after-tax IRR	31.3%

Each step is rounded to a tenth of a point, as the deck prints it.

More metrics


NPV at 10%	\$1,360,732
Year-one distribution	\$561,564
Multiple on invested capital	3.10×
Pre-tax IRR	36.27%
Installed cost	\$1,407,240

IRR build: the cash-only return, then the depreciation and credit effects. **More metrics:** NPV, the year-one distribution, the multiple and the pre-tax IRR.

Warnings & notes say what was assumed and what replaces it. Read them before you send the deck.

After-tax cash flow

Year 0 is the equity at close



[> Annual table](#)

After-tax cash flow by year, with the cumulative line. **Annual table** opens every year's figures.

Sensitivity

After-tax IRR

Case	After-tax IRR	Change
Base case	31.29%	—
Installed cost –10%	35.21%	+3.93 pts
Installed cost +10%	28.00%	–3.29 pts
GPU-hour prices –20%	24.00%	–7.29 pts
GPU-hour prices +20%	38.26%	+6.97 pts
Edge and on-demand utilisation –10 points	28.00%	–3.29 pts
Electricity rates +25%	31.11%	–0.18 pts
Charging sessions –25%	30.76%	–0.53 pts

Sensitivity: installed cost, GPU-hour prices, utilisation, electricity rates and charging sessions.

Below the sensitivity, **The operating schedule, year by year** shows each revenue and cost line over the term, **Deal comparison** sets the three structures side by side, and **Modelling conventions** lists the method.

The investor deck

The deck is drawn from the result in your company's brand: a cover, the one-page investor summary, the returns and tax basis, the annual cash flow, the assumptions and a closing page. **Print / save as PDF** prints it; **Download the annual CSV** exports the year-by-year figures.

INVESTOR ONE-PAGER · PREPARED 09 / 30 / 2026

Your Company

4 H200 pods + 12 EV ports in Austin, TX

The sponsor is seeking a \$1.41 million investment in a project that installs 192 H200 GPUs in 4 compute pods, 12 EV charging ports and a 100 kW / 200 kWh battery at Example apartments in Austin, TX, on the building's existing 316 kW service. The project owns the power and pod layer and receives 20% of gross compute revenue from the operator that owns the GPUs. An hour-by-hour model of the building, the chargers, the compute and the battery holds the service below 253 kW, and the owner captures the project's investment tax credit and depreciation benefits.

\$1.41M
total investment sought
100% equity, unlevered

31.3%
after-tax unlevered IRR
over the 10-year term

\$2,968,212
total investor returns
over life of project

~2.7 yrs
after-tax payback
cash + depreciation + ITC

HOW THE SYSTEM WORKS


192 H200 GPUs
4 pods, 158 kW serving

→


12 EV ports
Managed charging, 138 kW

→


100 kW / 200 kWh BESS
Holds the service limit and the peak

→


Example apartments
Multifamily (house meter), Austin

→


Existing service
316 kW at 480 V, shared

10-year analysis · 20% of gross compute revenue to the project

Your Company — Example: 4 GPU pods on a 400 A apartment service (illustrative) | Investor One-Pager

The investor one-pager, as printed.

Save scenario at the top keeps every input. Saved scenarios are listed under **Compare scenarios**, at the foot of the results, with their size, schedule, deal and returns side by side.

From Site Map: the Site Screen

In Site Map, open the **Compute** tab › **Size** › **Site Screen**. It reads what the drawing already knows: the ZIP, the service from the behind-the-meter inputs, the chargers placed, the battery and solar, the pods, and the Battery Sizer's bills. It screens the site on the server and prints ADVANCE, VERIFY or HOLD with the next gate. **Open the Compute Site Pro Forma** carries the site into this page, filled in.

Reading the numbers

Service limit	The service rating times the continuous loading limit, plus any upgrade. Everything is held under it.
Firm load	The building and the contracted compute (offtake, and edge contracts in their hours). It is never held back.
On-demand compute	GPU-hours sold as they are asked for. The only load the model holds back to keep the peak down.
Peak at the meter	The highest hour of the year the utility sees, after the battery and any held-back GPUs.
Every charger at full power	The nameplate check (NEC 625.42): all ports at full power plus the firm load. Above the limit, charging needs load management.
Not delivered	Charging energy a car could not get before it left.
After-tax IRR	Unlevered with 100% equity; the equity IRR with project debt.
Payback	After tax: cash, depreciation and the credit.

Good to know

- Planning figures are marked on the page and in the deck: the GPU-hour prices, the market's planning rate and the typical load. The operator's contract, the utility tariff and 12 months of bills or an 8760 replace them.
- A link can fill the page: **?zip=78701&s=400&pods=4&ports=12** (also kw, volts, portkw, bkw, bkwh, solar, annualKwh, hostType, gpu, pattern, name, city, host, street, utility).
- The model is one meter over a typical weather year with a constant PUE. Demand ratchets and solar export credit are not modelled.
- It is a screening model, not a compute operator's offer, a utility load study or a tax opinion.